THE UNIVERSITY of EDINBURGH informatics

Probabilistic Modelling and Reasoning
Tutorial 4 — Notes

Spring 2019
Zack Hodari

These notes are intended to give a summary of relevant concepts from the lectures which are helpful to complete the tutorial sheet. It is not intended to cover the lectures thoroughly. Learning this content is not a replacement for working through the lecture material and the tutorial sheet.

Factor graph — A factor graph represents an arbitrary function in terms of factors and their connections with variables. For example, a factor graph can represent a distribution written as a Gibbs distribution – $p(\mathbf{x}) = \frac{1}{Z}\prod_c \phi_c(\mathcal{X}_c)$ – where variables $x_i \in \mathbf{x}$ are represented with variable nodes (circles) and potentials ϕ_c are represented with factor nodes (squares). Edges connect each factor node ϕ_c to all its variable nodes $x_i \in \mathcal{X}_c$.

$$p(x_1, x_2, x_3, x_4) = \frac{1}{Z}\phi_1(x_1, x_2, x_3)\phi_2(x_3, x_4)\phi_3(x_4)$$

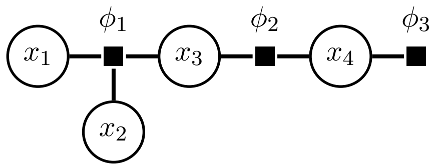

Variable elimination — Given a factorisation $p(\mathcal{X}) \propto \prod_c \phi_c(\mathcal{X}_c)$, we can compute the marginal $p(\mathcal{X} \setminus x^*)$ by exploiting the distributive property of the factorisation.

We sum out the variable x^* by first finding all factors $\phi_i(\mathcal{X}_i)$, such that $x^* \in \mathcal{X}_i$. We combine these factors into a new factor $\tilde{\phi}^*(\tilde{\mathcal{X}}^*) = \sum_{x^*}\prod_{i:x^*\in\mathcal{X}_i}\phi_i(\mathcal{X}_i)$. Note that $\tilde{\phi}^*$ does not depend on x^*, i.e. $\tilde{\mathcal{X}}^* = \bigcup_{i:x^*\in\mathcal{X}_i}(\mathcal{X}_i \setminus x^*)$. This is possible as products are commutative, and a sum can be distributed within a product as long as all terms depending on the variable(s) being summed come to the right of the sum.

$$p(\mathcal{X} \setminus x^*) \propto \sum_{x^*}\prod_c \phi_c(\mathcal{X}_c) \propto \left[\prod_{i:x^*\notin\mathcal{X}_i}\phi_i(\mathcal{X}_i)\right]\left[\sum_{x^*}\prod_{i:x^*\in\mathcal{X}_i}\phi_i(\mathcal{X}_i)\right] \tag{1}$$

$$\propto \left[\prod_{i:x^*\notin\mathcal{X}_i}\phi_i(\mathcal{X}_i)\right]\tilde{\phi}^*(\tilde{\mathcal{X}}^*) \tag{2}$$

When eliminating variables, order of elimination matters. However, optimal choice of elimination order is difficult. Picking variables greedily is a common heuristic, where the "best" x^* is the one that fewest factors ϕ_c depend upon.

Sum-product algorithm — Variable elimination (for factor trees) reformulated with "messages" which allows for re-use of computations already done. See table on following page.

Max-product algorithm — Same as the sum-product algorithm, but *max* replaces $\sum$.

$\mu_{\phi \to x}(x)$	Factor to variable $\mu_{\phi \to x}(x) = \sum_{x_1,\ldots,x_j} \phi(x_1,\ldots,x_j,x) \prod_{i=1}^{j} \mu_{x_i \to \phi}(x_i)$	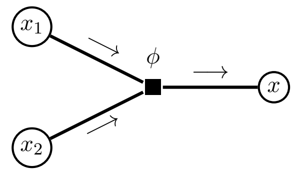
$\mu_{x \to \phi}(x)$	Variable to factor $\mu_{x \to \phi}(x) = \prod_{i=1}^{j} \mu_{\phi_i \to x}(x)$	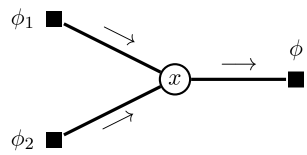
$\tilde{p}(x)$	Univariate marginals – unnormalised $p(x) \propto \prod_{i=1}^{j} \mu_{\phi_i \to x}(x)$	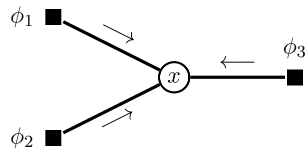
$\tilde{p}(x_1,\ldots,x_j)$	Joint marginals – unnormalised $p(x_1,\ldots,x_j) \propto \phi(x_1,\ldots,x_j) \prod_{i=1}^{j} \mu_{x_i \to \phi}(x_i)$